

Contents

1	Introduction	3
1.1	A Short Overview of Programming Languages	3
1.1.1	C, Strengths and Weaknesses	4
1.2	Shell Basics on Linux	5
2	Fundamentals of C Language	9
2.1	Writing a Simple Program, <code>Hello, McMaster!</code>	9
2.1.1	A General Form of a Program	9
2.1.2	Compiling and Executing the Program	11
2.2	Integer Data Type.....	13
2.2.1	Binary Representation	13
2.2.2	Using <code>printf</code> and <code>limits.h</code> to Get the Limits of Integers	15
2.2.3	Declaring and Naming with and without Initializing.....	19
2.2.4	Constant Variables	20
2.2.5	Arithmetic Operations on Integers.....	21
2.2.6	A Simple Example for Integer Overflows.....	22
2.2.7	Fixed-width Integer types	25
2.3	Characters and strings	27
2.4	Floating-point Numbers	29
2.4.1	Rounding Error in Floating-point Numbers	33
2.4.2	Type Conversion.....	34
2.4.3	Arithmetic Operations on Floating-point Numbers.....	36
2.5	Mathematical Functions.....	40
2.6	Statements	42
2.6.1	Comparison: <code>if</code> and <code>switch</code>	42

2.6.2	Loops and Iterations: <code>while</code>	49
2.6.3	Loops and Iterations: <code>do</code>	50
2.6.4	Loops and Iterations: <code>for</code>	51
2.6.5	Nested loops	52
2.6.6	Loops and Iterations: Controlling the Loop	54
2.6.7	Variable Scope	56
2.7	Arrays	56
2.8	Functions	56
3	Intermediate Topics in C	56
4	Advanced Topics in C	57
5	A Short Overview of ANN	57

1 Introduction

1.1 A Short Overview of Programming Languages

Programming languages can be categorized based on their level of abstraction, which refers to how closely they resemble the underlying hardware operations of a computer. At the lowest level, we have **machine code**, which consists of binary instructions that are directly executed by the computer's hardware. Above machine code, we have **assembly language**, which uses **human-readable** characters to represent the low-level instructions. However, assembly language is specific to each CPU architecture, so it can differ between different types of processors.

Moving up the hierarchy, we encounter **C**, which was created in around 1970 as a by-product of UNIX based operating systems. **C** is considered a slightly higher-level language compared to **assembly language**. It provides a set of abstract statements and constructs that are closer to human-readable text. **C** code can be compiled using a program called **compiler**, which translates the **C** code into **machine code** that can be executed on any CPU. This portability is one of the reasons why **C** is often referred to as "**portable assembly**." **C** allows programmers to efficiently write code with a relatively low level of abstraction.

Above **C**, we find **C++**, which was created around 1985. **C++** is an extension of **C** and introduces **object-oriented programming** concepts. It includes the ability to define **classes**, which encapsulate data and behavior into reusable structures. However, for the purpose of this course, we won't delve into the specifics of object-oriented programming.

Moving further up the ladder, we have **Java** and **C#**, which are considered **mid-level** languages. These languages restrict certain low-level operations available in **C** and **C++**, for example, managing memory environments. Instead of allowing direct memory allocation, **Java** and **C#** handle memory management themselves, offering automatic memory allocation and garbage collection. This trade-off provides programmers with increased security and simplifies memory management, but it also limits some of the flexibility and control offered by lower-level languages.

At the **highest level**, we have interpreted languages like **Python** and **JavaScript**. These languages are considered highly abstracted and provide significant levels of convenience and ease of use for developers. Interpreted languages do not require a separate compilation step; instead, they use an **interpreter** to execute the source code directly. The interpreter reads the code **line-by-line**, executing each instruction as it encounters it. This line-by-line execution allows for more dynamic and interactive programming experiences but can result in slower performance compared to compiled languages.

1.1.1 C, Strengths and Weaknesses

Weaknesses

While C has many strengths, it also has a few weaknesses that developers should be aware of:

1. **Error-prone:** Due to its flexibility, C programs can be prone to errors that may not be easily detectable by the compiler. For example, missing a semicolon or adding an extra one can lead to unexpected behavior, such as infinite loops! It means you are waiting for hours to get the result, then you figure out it shouldn't take that much time! Don't worry there are some ways to avoid it!

2. **Difficulty in Understanding:** C can be more difficult to understand compared to higher-level languages. It requires a solid understanding of low-level concepts, such as memory management and pointers. The syntax and usage of certain features, like pointers and complex memory operations, may be unfamiliar to beginners or programmers coming from higher-level languages. This learning curve can make it more challenging for individuals new to programming to grasp and write code in C effectively.

3. **Limited Modularity:** C lacks built-in features for modular programming, making it harder to divide large programs into smaller, more manageable pieces. Other languages often provide mechanisms like namespaces, modules, or classes that facilitate the organization and separation of code into logical components. Without these features, developers must rely on manual techniques, such as using header files and carefully organizing code structure, to achieve modularity in C programs. This can lead to codebases that are harder to maintain and modify as the project grows in complexity.

Strength

C programming language possesses several strengths that have contributed to its enduring popularity and wide-ranging applicability:

1. **Efficiency:** C is renowned for its efficiency, allowing programs written in C to execute quickly and consume minimal memory resources. It provides low-level access to memory and hardware, enabling developers to optimize their code for performance-critical applications.

2. **Power:** C offers a rich set of data types and a flexible syntax, enabling developers to express complex algorithms and manipulate data efficiently. Its extensive standard library provides numerous functions for tasks such as file I/O, memory management, and string manipulation, allowing programmers to accomplish a lot with concise and readable code.

3. **Flexibility:** The versatility of C is evident in its widespread use across different domains and industries. C has been employed in diverse applications, including embedded systems, oper-

ating systems, game development, scientific research, commercial data processing, and more. Its flexibility makes it a suitable choice for a broad range of programming tasks.

4. Portability: C is highly portable, meaning that programs written in C can be compiled and run on a wide range of computer systems, from personal computers to supercomputers. The C language itself is platform-independent, and compilers are available for various operating systems and architectures.

5. UNIX Integration: C has deep integration with the UNIX operating system, which includes Linux. This integration allows C programs to interact seamlessly with the underlying system, making it a favored language for system-level programming and development on UNIX-based platforms.

1.2 Shell Basics on Linux

To create a new folder (directory) in Linux using the terminal, you can use the `mkdir` command. Here's how you can do it:

Open your terminal application. This can vary depending on the Linux distribution you're using. You can typically find the terminal in Applications menu or by pressing `Ctrl+Alt+T`. Navigate to the location where you want to create the new folder. You can use the `cd` command to change directories. For example, if you want to create the folder in your home directory, you can use `cd ~` to navigate there. Once you're in the desired location, use the `mkdir` command followed by the name you want to give to the new folder. For example, to create a folder called "COMPSCI1XC3," you would type `mkdir COMPSCI1XC3`. Press `Enter` to execute the command.

In the world of Linux commands hold the power to shape the digital landscape. Let's do start with some basic commands:

1. Current directory: To know your current directory, use the command `pwd`. It will display the path of the directory you are currently in. To view all the files and folders in the current directory, use the command `ls`. This will provide a list of all the files and folders. Look for a folder named `COMPSCI1XC3` in the list of files and folders (if you have did the previous step successfully!). Once you find it, you'll use the following steps to navigate and work within that folder. To open the `COMPSCI1XC3` folder in the file manager, use the command `open COMPSCI1XC3`.

2. Moving back and forth in directories: To change your current directory to the `COMPSCI1XC3` folder, use the command `cd COMPSCI1XC3` in the terminal. This command allows you to move into the specified folder (use `pwd` to double check!) If you want to go back to the initial directory, simply use the command `cd` without any arguments. This will take you back to the `/home/username` directory, where `username` is what you picked during installation

of Linux. To go back one folder, use the command `cd ..`. This will navigate you up one level in the directory structure.

3. Making a file: Now, go back to `COMPSCI1XC3` directory. If you want to create a text file named "AnneMarie", use the command `nano AnneMarie.txt`. This will open the file editor where you can write and edit text. To open the `AnneMarie.txt` file in the Nano text editor, use the command `nano AnneMarie.txt` again. This allows you to make changes to the file. If you wish to save your changes, press `Ctrl+O` and then press `Enter`. This will save the file. Alternatively, if you want to save and exit the Nano text editor, press `Ctrl+X` and then press `Enter`. This will save the changes and return you to the terminal.

To reopen the `AnneMarie.txt` file in the terminal using the Nano text editor, use the command `nano AnneMarie.txt`. If you prefer to open the `AnneMarie.txt` file in the Windows environment, use the command `open AnneMarie.txt` after saving the file. This will open the file using the default application associated with `.txt` files.

4. Copy and Paste: To copy the `AnneMarie.txt` file and paste it into another directory, you can use the `cp` command, like:

```
cp AnneMarie.txt /path/to/destination/directory/
```

Replace `/path/to/destination/directory/` with the actual path of the directory where you want to paste the file.

5. View a file: The `cat` command is used to display the contents of a file in the terminal. It can be helpful when you want to quickly view the contents of files. Here's an `.txt` file example: `cat AnneMarie.txt`. Running this command will print the contents of `AnneMarie.txt` in the terminal window. Easier way just double click on the file!

6. Finding a folder of file: To find a file or folder using the `locate` command or similar commands, you need to have the appropriate indexing database set up on your system. The `locate` command searches this database for file and folder names. For example, `locate AnneMarie.txt`, will search the indexing database for any files or folders with the name `AnneMarie.txt` and display their paths if found. In this case, nothing is shown in terminal. Since we just made this file, the database including directories is not updated to include this file. To update the directory database in all memory, we have to use the command `updatedb`, and you will probably get the following message:

```
/var/lib/plocate/: Permission denied
```

indicating that you do not have the necessary permissions to access the directory. To solve the issue and update the directories index successfully, use the `sudo updatedb` which prompts you to a password which you picked. Enter the password, it might not be shown when you are entering

the password, then press `Enter`. By this command, you are running the `updatedb` with superuser (root) privileges. `sudo` stands for "Super User Do" and is a command that allows regular users to execute commands with the security privileges of the superuser. Now you can use command `locate`, but this time it shows the directory that the file is saved.

Alternatively, you can use:

```
find /path/to/search/directory -name AnneMarie.txt
```

Replace `/path/to/search/directory` with the directory where you want to start the search.

Tips! Debugging is an essential aspect of programming. No programmer possesses complete knowledge or can remember every command, especially when working with multiple programming languages. However, it is crucial to know where to find the answers. Here are two approaches to tackle a specific issue:

1. Use internet. For instance, search for "permission denied on Linux." Take a quick look at the top search results, paying particular attention to reputable sources such as [Stack Overflow](#). Spend a few minutes scrolling through the search results.
2. Engage with ChatGPT by asking a specific question related to the problem you encountered. For example, you could ask, "I tried `updatedb` on Linux terminal, and it gives me Permission denied. How can I solve it?"

7. Remove: The `rm` command is used to remove (delete) files and directories. It's important to exercise caution when using this command, as deleted files cannot be easily recovered. Here's an example to remove the `AnneMarie.txt` file: `rm AnneMarie.txt`. This command will permanently delete the `AnneMarie.txt` file. Be sure to double-check the file name and verify that you want to delete it.

The `rmdir` command is used to remove (delete) empty directories. It cannot remove directories that have any files or subdirectories within them. Here's an example: `rmdir empty_directory`. Replace `empty_directory` with the name of the directory you want to remove. This command will only work if the directory is empty. If there are any files or subdirectories inside it, you'll need to delete them first or use the `rm` command with appropriate options to remove them recursively.

8. Let's checkout some OS characteristics.

Linux distribution: distThe `lsb_release -a` command provides comprehensive information about your Linux distribution, including the release version, codename, distributor ID, and other relevant details. The `-a` option is a command-line option or flag that stands for "all." When used with the `lsb_release` command, it instructs the command to display all available information

about the Linux distribution. It is a convenient way to quickly check the specifics of your Linux distribution from the command line. My Linux distribution is:

```
Distributor ID: Ubuntu
Description: Ubuntu 22.04.2 LTS
Release: 22.04
Codename: jammy
```

CPU: The `lscpu` command is used to gather and display information about the CPU (Central Processing Unit) and its architecture on a Linux system. It provides detailed information about the processor, including its model, architecture, number of cores, clock speed, cache sizes, and other relevant details. **Some** of details for my system:

```
Architecture:          x86_64
  CPU op-mode(s):      32-bit, 64-bit
  Address sizes:       39 bits physical, 48 bits virtual
  Byte Order:          Little Endian
CPU(s):                8
  On-line CPU(s) list: 0-7
Vendor ID:             GenuineIntel
  Model name:          Intel(R) Core(TM) i7-7700HQ CPU @ 2.80GHz
  CPU family:          6
  Model:               158
  Thread(s) per core:  2
```

Disk space available: The `df -h` command is used to display information about the disk space usage on your Linux system. In this command, `df` stands for **disk free**, and `-h` is a command-line option or flag that stands for **human-readable**. The

```
Filesystem      Size  Used Avail Use% Mounted on
/dev/sda2       228G  113G   104G   53% /
```

`/dev/sda2` is a partition identifier that represents the second partition on the first SATA (or SCSI) disk (sda). It typically contains the main file system of your Linux system. The actual partition identifiers may differ depending on your system's configuration and the number of disks or partitions present.

Information about the RAM: In the command `free -h`, `free` represents the command used to display memory usage statistics, and `-h` is a command-line option that stands for **human-readable**.

```
total      used      free      shared  buff/cache   available
```


Mem:	15Gi	4.4Gi	3.3Gi	637Mi	7.8Gi	10Gi
Swap:	2.0Gi	0B	2.0Gi			

The **Swap** space is a portion of the hard drive that is used as virtual memory by the operating system. It acts as an extension to the physical memory (RAM) and allows the system to temporarily store data that doesn't fit into the RAM.

9. Short keys in terminal: `Ctrl+L` to remove the history of terminal. Pressing up key takes you to previous commands. To copy from terminal you have to press `Cntrl + Shift + C` and to paste a command press `Ctrl + shift + V`.

2 Fundamentals of C Language

This section is designed to provide an introduction to fundamental concepts in C, including data types, variables, and control flow statements such as if and loops. Some of students with prior knowledge of programming might already know these concepts, but we will delve into each topic extensively and provide detailed explanations. Our goal is to ensure that everyone, regardless of their programming background, can grasp these essential concepts effectively. We will illustrate each concept with practical examples and problem-solving exercises to enhance understanding, as these concepts form the backbone of C programming.

2.1 Writing a Simple Program, `Hello, McMaster!`

2.1.1 A General Form of a Program

Let's make a new file using `nano Hello.c` where the `.c` extension represent a C file. Before closing the the opened file in terminal copy and paste the following code in the file and save it. To paste the the copied section you have to use `Ctrl + shift + V`. Press `Ctrl + X` then `y` and press Enter to save before exit.

```
// this code is written by Pedram
#include <stdio.h>

// the main function
int main(void) {
    /* calling "printf" function
    the "defination" is included in "stdio" library */
    printf("Hello, McMaster!\n");
}
```

```
}
```

Open the file by `open Hello.c`. This is a more user friendly environment to edit the code. Maybe the main difference is I can click different parts of the code and edit the code. During you lab lectures this week, you TAs will discuss Visual Studio Code IDE (Integrated Development Environment) which support many programming languages, including C, as well as many useful tools such as automatic code formatting. I strongly suggest to install it and set it up for C format for the next lectures. If I want to modify the code in terms of extra spaces in the code using `open Hello.c` it will take me some time. But if you have Visual Studio Code installed, in the terminal enter `code Hello.c` to open the same code in Visual Studio Code (VSC). Right click on the code and choose **Format Document** which automatically modifies the code in C format.

Tips! Get your hands dirty! I promise you that by only reading these notes you will NOT be able to learn programming. You have to write, modify, and change every code in these notes. Play with them like you play video games!

Anything starts with two slashes, `//`, until the end of the line is a comment and it will not be executed. To make multiple lines as comments, you can place them between `/*` and `*/`. The following lines in the code are all comments:

```
// this code is written by Pedram

// the main function

/* calling "printf" function
the "defination" is included in "stdio" library */
```

Based on the IDE that you are using the colour of the comments might be different. I have said earlier that understanding C codes especially written by someone else might be difficult. Comments are supposed to be helpful to remember or tell others what is the purpose of each part of the code or even how it works.

So the first line which is really executed is `include <stdio.h>`. Usually at the beginning of a code, we include libraries used in the code. `include <stdio.h>` tells C to read **header file** named `stdio` with extension `.h` which stands for **header file**. In this library the definition of function `printf` is mentioned. Any time you forget what library a function is in, you can ask your best friend **Google** or **ChatGPT** by simply asking "what library in C printf is in."

The second line that will be executed is `int main(void)` as a function. In C, defining a function starts with the type of value the function returns, which in this case it is `int`. The `int` means **integer** and the function `main` will automatically return the value `0` if there is no error in

the code.

The `main` is the name of function. This function must be included in all C programs, and you have to write your code inside this function.

Between the parentheses, `void` what inputs your program needs and requires the user to give the input(s) every time the code is going to be executed. In this case, we don't have any input so it is `void`, or we can define `int main()` instead of `int main(void)`. Both means the program requires no input(s).

Anything between `{`, right after `int main(void)`, and `}`, at the end of the code, is called **compound statement** or **code block**, which belongs to `main` function. So the following statement is the general format that we have to include in all C programs:

```

Including Library

int main(){

    your statements

}
```

In this code, `printf("Hello, McMaster!\n")` is the only statement we have and `printf` function is used to print a text. Between `(` and `)` the arguments for the function `printf` is given, which is a string or text to be printed!

A string is sequence of characters enclosed within quotation marks. `"Hello, McMaster!\n"` is the string that will be printed. The backslash `\` before `n` tells C to go to (n)ext line after printing `Hello, McMaster!`. Try to compile and run the code two times with `\n` and two times without `\n` and see the difference in the terminal! At the end of this line there is `;` indicating the command for this line has ended.

A reminder, the closing bracket `}` at the end, means the end of function `main`!

2.1.2 Compiling and Executing the Program

Now we have to convert the program, the file `Hello.c`, to a format that machine can understand and execute. It usually involves, **Pre-processing**, **Compiling**, and **Linking**. During **Pre-processing** some modification will be done automatically to the code, before making the executable **object** code in **Compiling** step. These two steps are done at the same time and automatically, so won't need to be worried about it! In the Linking step, the library included will be linked to the executable file. Since this code is using a standard library, `<stdio.h>`, it happens

automatically. Again, you don't need to be worried!!!

To compile the program in UNIX based OS, usually `cc` is used. Open a terminal in the directory where `Hello.c` file is located (please refer to the section Shell Basics on Linux). Type `cc Hello.c`. Check out the directory. A new executable **object** file named `a.out` by default will be created.

If you see the following error after executing `cc Hello.c`, it means you are not in the same directory where `Hello.c` is located. A reminder, use `pwd` to check your current directory.

```
cc1: fatal error: Hello.c: No such file or directory
compilation terminated.
```

Visual Studio Code! You may open the code in Visual Studio Code (VSCode) by running `code Hello.c` in the terminal. Again make sure current directory is where `Hello.c` is located, otherwise you will see an error indicating the is no such a file in this directory. At the top of VSCode opened, press **View** and click **Terminal**. A window will be opened at the bottom, named **TERMINAL**. There is no difference between this Terminal and terminal you open using `Ctrl + Alt + T`. Make sure the directory of this Terminal in VScode is the where `Hello.c` is located to avoid the error I have mention before running `cc Hello.c` in the terminal.

After `cc Hello.c` if there is no problem, you should be able to see a new executable **object** file named `a.out` by default in the directory.

At this stage the code is compiled, the **object** `a.out` readable by machine is created. This means the program is translated to a language that machine can understand and it saved in the **object** file `a.out`. Now it is time to tell the machine to run the translated code and show us the result. Run `./a.out` in the terminal. Again, make sure you are in the directory where `a.out` is located. It is done! you should be able to see the result!

Hello, McMaster!

I have mentioned the name `a.out` is given to the **object** file by default. I can define any name I want. Lets remove the previous object file by `rm a.out` and make a new one with the name "Brad" by using `cc -o Brad Hello.c`. The compiler `cc` has many options, and `-o` means translate the program `Hello.c` to the machine language with an (o)bject file named `Brad`. Now you must be able to see an **object** file named `Brad` in the same directory. If you run the command line `./Brad` in the same directory, you should be able to see the output.

GCC compiler Another popular compiler in C is GCC (GNU Compiler Collection) compiler supported by Unix OS. It is known for its robustness, efficiency, and compatibility with multiple platforms. GCC supports various optimizations and provides comprehensive error checking, making it a reliable choice for compiling C code.

One notable feature of GCC is its similarity to the "cc" compiler command. This similarity stems from the fact that on many Unix-like systems, the "cc" command is often a symbolic link or an alias for GCC. Therefore, using "cc" to compile your code essentially invokes the GCC compiler with its default settings. GCC offers numerous options to control various aspects of the compilation process, such as optimization levels, debugging symbols, and specific target architectures.

You can compile the same program using GCC instead of `cc` by executing `gcc -o Brad Hello.c` in the terminal.

The **object** file `./Brad` is executable in any Unix based OS. It is software of application you developed!

2.2 Integer Data Type

In contemporary C compilers, there is support for a range of integer sizes spanning from 8 to 64 bits. However, it is worth noting that the specific names assigned to each size of integer may differ among different compilers, leading to potential confusion among developers. First we need to know how data is stored in computers.

2.2.1 Binary Representation

In the world of computers, information is stored and processed as sequences of bits, representing either a 0 or a 1. This fundamental concept forms the basis of how data is handled by computer systems.

Consider the scenario where we have a fixed number of bits, denoted by N , to represent our numbers. Let's take $N=4$ to save an `int` number where `int` in C stands for **integer** number. In this case, we are allocating nine bits to express our numerical values. If the integer is `unsigned` then:

- $0 = 0000$
- $1 = 0001$
- $2 = 0010$
- $3 = 0011$

- $4 = 0100$
- $5 = 0101$
- $6 = 0110$
- $7 = 0111$
- $8 = 1000$
- $9 = 1001$
- $10 = 1010$
- $11 = 1011$
- $12 = 1100$
- $13 = 1101$
- $14 = 1110$
- $15 = 1111 = 2^N - 1$

More example?! If $N = 3$ then for `unsigned` integer we would have:

- $0 = 000$
- $1 = 001$
- $2 = 010$
- $3 = 011$
- $4 = 100$
- $5 = 101$
- $6 = 110$
- $7 = 111 = 2^N - 1$

which means the range of numbers can be saved as `unsigned` integer in C is from 0 to $2^N - 1$. For `signed` integers, the negative numbers can be obtained by inverting the bits then adding one to the result. This method is called the two's complement operation.

- $-8 = 1000 = -2^{N-1}$
- $-7 = (\text{inverting } 0111, \text{ we have } 1000, +1 \text{ is:})1001$
- $-6 = (\text{inverting } 0110, \text{ we have } 1001, +1 \text{ is:})1010$
- $-5 = (\text{inverting } 0101, \text{ we have } 1010, +1 \text{ is:})1011$
- $-4 = (\text{inverting } 0100, \text{ we have } 1011, +1 \text{ is:})1100$
- $-3 = (\text{inverting } 0011, \text{ we have } 1100, +1 \text{ is:})1101$
- $-2 = (\text{inverting } 0010, \text{ we have } 1101, +1 \text{ is:})1110$
- $-1 = (\text{inverting } 0001, \text{ we have } 1110, +1 \text{ is:})1111$
- $0 = 0000$
- $1 = 0001$
- $2 = 0010$

- $3 = 0011$
- $4 = 0100$
- $5 = 0101$
- $6 = 0110$
- $7 = 0111 = 2^{N-1} - 1$

If you have noticed, you can see the 1000 is signed to -8. In total with 4 bits, machine can have a combination of 2^4 zero and ones. It can be seen that all positive numbers starts with 0, and negative ones start with one. Therefore, it is accepted to sign 1000 bit sequence to the lowest number -8.

More example?! If $N = 3$ then for `unsigned` integer we would have:

- $-4 = 100 = -2^{N-1}$
- $-3 = (\text{inverting } 011, \text{ we have } 100, +1 \text{ is:})101$
- $-2 = (\text{inverting } 010, \text{ we have } 101, +1 \text{ is:})110$
- $-1 = (\text{inverting } 001, \text{ we have } 110, +1 \text{ is:})111$
- $0 = 000$
- $1 = 001$
- $2 = 010$
- $3 = 011 = 2^{N-1} - 1$

which means the range of numbers can be saved as `signed` integer in C is from -2^{N-1} to $2^{N-1} - 1$. Exceeding this range can cause errors, probably unseen, which it is called **integer overflow**.

2.2.2 Using `printf` and `limits.h` to Get the Limits of Integers

Let's see some of these limits using `limits.h` library. Make a new file using `nano limit.c` which `limit` is the name of program given by you, and `.c` is the extension which stands for C files. Copy and past the following code. Press `Ctrl + X` then `y` and Enter to save and close the file. Open the code in VScode, if you have it as IDE, by entering `code limit.c` in your terminal. Make sure you are in the same directory where you made this file.

```
// This code is written by ChatGPT
#include <stdio.h>
#include <limits.h>

int main() {
```

```

printf("Size of char: %zu bits\n", 8 * sizeof(char));
printf("Signed char range: %d to %d\n", SCHAR_MIN, SCHAR_MAX);
printf("Unsigned char range: %u to %u\n", 0, UCHAR_MAX);
printf("\n");

printf("Size of int: %zu bits\n", 8 * sizeof(int));
printf("Signed int range: %d to %d\n", INT_MIN, INT_MAX);
printf("Unsigned int range: %u to %u\n", 0, UINT_MAX);
printf("\n");

printf("Size of short: %zu bits\n", 8 * sizeof(short));
printf("Signed short range: %d to %d\n", SHRT_MIN, SHRT_MAX);
printf("Unsigned short range: %u to %u\n", 0, USHRT_MAX);
printf("\n");

printf("Size of long: %zu bits\n", 8 * sizeof(long));
printf("Signed long range: %ld to %ld\n", LONG_MIN, LONG_MAX);
printf("Unsigned long range: %u to %lu\n", 0, ULONG_MAX); //
    change %u to lu
printf("\n");

printf("Size of long long: %zu bits\n", 8 * sizeof(long long));
printf("Signed long long range: %lld to %lld\n", LLONG_MIN,
    LLONG_MAX);
printf("Unsigned long long range: %u to %llu\n", 0, ULLONG_MAX);
    // change %u to llu
}

```

Let's check the code line-by-line.

1. `#include <limits.h>`: This is a preprocessor directive that includes the header file `limits.h` in the C program. The `limits.h` header provides constants and limits for various data types in the C language, such as the minimum and maximum values that can be represented by different types.

2. Placeholders: is a special character or sequence of characters that is used within a formatted string to represent a value that will be substituted during runtime. Placeholders are typically used in functions like `printf` or `sprintf` to dynamically insert values into a formatted output. When the program runs, the placeholders are replaced with the actual values passed as arguments to the formatting function. The values are appropriately converted to match the format specifier

specified by the corresponding placeholders.

Placeholders provide a flexible way to format output by allowing dynamic insertion of values. They help in producing formatted and readable output based on the specified format specifiers and the provided values.

- `%zu`: This is a placeholder used with `printf` to print the value of an `unsigned integer`.
- `%d`: This is a placeholder used to `printf` the value of a `signed integer`.
- `%u`: This is a placeholder used to `printf` the value of an `unsigned integer`.
- `%ld`: This is a placeholder used to `printf` the value of a `signed long integer`.
- `%lu`: This is a placeholder used to `printf` the value of an `unsigned long integer`.
- `%lld`: This is a placeholder used to `printf` the value of a `signed long long integer`.
- `%llu`: This is a placeholder used to `printf` the value of an `unsigned long long integer`.

3. `sizeof()`: This is an operator in C that returns the size of a variable or a data type in bytes. In the given code, `sizeof()` is used to determine the size of different data types (e.g., `char`, `int`, `long`, `long long`). The result of `sizeof()` is then multiplied by 8 to obtain the size in bits.

4. `printf("\n")`: This line of code is using `printf` to print a newline character `\n`. It adds a line break, resulting in a new line being displayed in the console output.

5. `NAME_MIN` and `NAME_MAX`: The code refers to variables like `SCHAR_MIN`, `SCHAR_MAX`, `INT_MIN`, `INT_MAX`, etc. These variables are predefined in the C library, specifically in the `limits.h` header file. They represent the minimum and maximum values that can be stored in the respective data types (e.g., `char`, `int`, `short`, `long`, `long long`).

6. `char`, `int`, `short`, `long`, and `long long`: These are data types in the C language. They represent different ranges of integer values that can be stored. The code provided displays the size and range of each of these data types, both signed and unsigned.

This time we compile the code with more options `gcc -Wall -W -std=c99 -o limit limit.c`, where:

`-Wall`: This flag enables a set of warning options, known as "all warnings." It instructs the compiler to enable a comprehensive set of warning messages during the compilation process. These warnings help identify potential issues in the code, such as uninitialized variables, unused variables, type mismatches, and other common programming mistakes. By enabling `-Wall`, you can ensure that a wide range of warnings is reported, assisting in the production of cleaner and more reliable code.

`-W`: This flag is used to enable additional warning options beyond those covered by `-Wall`. It allows you to specify specific warning options individually. Without any specific options following `-W`, it enables a set of commonly used warnings similar to `-Wall`. By using `-W`, you have more control over the warning messages generated by the compiler.

`-std=c99`: This flag sets the C language standard that the compiler should adhere to. In this case, c99 indicates the **C99 standard**. The **C99 standard** refers to the ISO/IEC 9899:1999 standard for the C programming language. It introduces several new features and improvements compared to earlier versions of the C standard, such as support for variable declarations anywhere in a block, support for `//` single-line comments, and new data types like `long long`. By specifying `-std=c99`, you ensure that the compiler follows the **C99 standard** while compiling your code.

`-o limit`: This flag is used to specify the output file name. In this case, it sets the output file name as "limit". The compiled binary or executable will be named "limit" as a result.

`limit.c`: This is the source file that contains the C code to be compiled.

After compiling the code, I can run the object file `limit` in the directory by `./limit`. This is the result I get in **my computer**!

```
Size of char: 8 bits
Signed char range: -128 to 127
Unsigned char range: 0 to 255

Size of int: 32 bits
Signed int range: -2147483648 to 2147483647
Unsigned int range: 0 to 4294967295

Size of short: 16 bits
Signed short range: -32768 to 32767
Unsigned short range: 0 to 65535

Size of long: 64 bits
Signed long range: -9223372036854775808 to 9223372036854775807
Unsigned long range: 0 to 18446744073709551615

Size of long long: 64 bits
Signed long long range: -9223372036854775808 to 9223372036854775807
Unsigned long long range: 0 to 18446744073709551615
```

We will find out about the importance of knowing limits in section [A Simple Example for](#)

Integer Overflows.

2.2.3 Declaring and Naming with and without Initializing

To declare a variable you can simply:

```
int Pedram;
```

where `int` is the type of variable and `Pedram` is the name of variable. You can after this line of code calculate the value of `Pedram`. You may also declare the value for this variable when it is initialized:

```
int Pedram = 10;
```

Tips! There are some reserved names that you cannot use for your variables, including: `auto`, `break`, `case`, `char`, `const`, `continue`, `default`, `do`, `double`, `else`, `enum`, `extern`, `float`, `for`, `goto`, `if`, `inline`, `int`, `long`, `register`, `restrict`, `return`, `short`, `signed`, `sizeof`, `static`, `struct`, `switch`, `typedef`, `union`, `unsigned`, `void`, `volatile`, `while`, `_Bool`, `_Complex`, `_Imaginary`, `#define`, `#include`, `#undef`, `#ifdef`, `#endif`, `#ifndef`, `#if`, `#else`, `#elif`, `#pragma`, and more! Don't worry if you use these you will see some errors and warnings when compiling the the code!

More Tips! About naming style, if you use only characters like `a,b,c, ..., z`, no one can follow your code or what is the purpose of this variable. So:

- Use meaningful and descriptive names that convey the purpose or nature of the variable. for example, `rectangle_height` and `triangle_width`
- Avoid excessively long names, but provide enough clarity to understand the purpose of the variable.
- Follow a consistent naming convention, such as `camelCase` or `snake_case`.
- use comments if necessary to explain the purpose or usage of a variable.

If you are compiling the code you can use `-Wextra` flag for uninitialized variables. Let's try the following code with and without.

```
#include <stdio.h>

int main() {
```

```
int x;
int y = x + 5; // Using uninitialized variable x
printf("%d\n", y);
}
```

Compiling the code with `gcc -o Pedram Pedram.c`, then executing the program with `./Pedram`. I get the following result:

```
-1240674203
```

If I execute the code one more time, `./Pedram`, without even compiling the code, I get:

```
233918565
```

What is going on???? When you run this code, you may get different output each time because the value of `x` is unspecified and can contain any arbitrary value. The variable `x` could be storing whatever value was previously in that memory location, and performing calculations with such a value can lead to unexpected results.

Let's compile the code but this time with `gcc -Wextra -o Pedram Pedram.c`. In my terminal it tells me `Using uninitialized variable x` and mentioning the line of code this issue is happening. At this point, the source code `Pedram.c` is compiled and the `Pedram` object is available. It means I can execute the program, but I am aware of the fact that this will result a wrong and an unexpected result. There are some warning you have to take serious even more than error! I could see the same warning by compiling the code using `gcc -Wall -W -std=c99 -o Pedram Pedram.c`

2.2.4 Constant Variables

Constant variables are declared using the `const` keyword, indicating that their value cannot be modified once assigned. Here's an example:

```
#include <stdio.h>

int main() {
    const int MAX_VALUE = 100;

    printf("Max value: %d\n", MAX_VALUE);
}
```

Start developing this habit to mention the constant values during programming which make you code to be more understandable. After using `const` for variable `MAX_VALUE`, I cannot change

the value for this variable. You can try to do it and you will see errors like:

```
expression must be a modifiable lvalue
```

or

```
assignment of read-only variable 'MAX_VALUE'
```

2.2.5 Arithmetic Operations on Integers

Let's play with some of these integer numbers using arithmetic operations.

```
#include <stdio.h>

int main() {
    int a = 5;
    int b = 3;
    int sum = a + b;
    int difference = a - b;
    int product = a * b;
    int quotient = a / b;
    int remainder = a % b;

    printf("Sum: %d\n", sum);
    printf("Difference: %d\n", difference);
    printf("Product: %d\n", product);
    printf("Quotient: %d\n", quotient);
    printf("Remainder: %d\n", remainder);

    return 0;
}
```

after compiling the code and executing the object file, you should get the following results:

```
Sum: 8
Difference: 2
Product: 15
Quotient: 1
Remainder: 2
```

2.2.6 A Simple Example for Integer Overflows

Why we tried to understand the limits of integer variables? If you define an `int` value equal to $2,147,483,647 + 1$, you will encounter a phenomenon known as **integer overflow**. In C, when an arithmetic operation results in a value that exceeds the maximum representable value for a given integer type, the behavior is undefined.

In most cases, when an integer overflow occurs, the value will "wrap around" and behave as if it has rolled over to the minimum representable value for that integer type. In the case of a 32-bit `int`, which has a maximum value of $2,147,483,647$, adding 1 to it will result in an **integer overflow**.

The exact behavior after the overflow is undefined, meaning it's not guaranteed what will happen. However, it is common for the value to wrap around to the minimum value for the `int` data type, which is typically $-2,147,483,648$ for a 32-bit signed integer. Let's try an example:

```
#include <stdio.h>
#include <limits.h>

int main() {
    int value = INT_MAX + 1;
    printf("Value: %d\n", value);

    return 0;
}
```

Compile the code using `gcc -o Anna Anna.c` where `Anna.c` is the source code, and `Anna` is the object file. Run the code with `./Anna`. The result I get in **my machine** is:

```
Value: -2147483648
```

while I was expecting to see $2,147,483,648$.

Let's see another example. Make a new file using terminal by `nano Pedram.c`. Copy and paste the following code using `Ctrl + Shift + V`. Press `Ctrl + X`, then press `y` and Enter to save the changes made to new file named `Pedram` with extension `.c`. In the following code the limits for each type is given in comments, which these limits are based on **my machine** and it might be different in yours. We found these limits by running the code in section **Using `printf` and `limits.h` to Get the Limits of Integers**.

```
#include <stdio.h>

int main() {
```

```

// ----- char -----
char Ped_RealChar = 'P';
// char range: -128 to 127
char Ped_NumChar = 80;
// unsigned char range: 0 to 255
unsigned char Ped_NumChar_unsigned = 252;
// ----- short -----
// short range: -32768 to 32767
short Ped_sh = -1234;
// unsigned short range: 0 to 65535
unsigned short Ped_sh_unsigned = 56789;
// ----- int -----
// int range: -2147483648 to 2147483647
int Ped_int = -42;
// unsigned int range: 0 to 4294967295
unsigned int Ped_int_unsigned = 123456;
// ----- long -----
// long range: -9223372036854775808 to 9223372036854775807
long Ped_long = -9876543210;
// unsigned long range: 0 to 18446744073709551615
unsigned long Ped_long_unsigned = 9876543210;
// # ----- long long -----
// long long range: -9223372036854775808 to
// 9223372036854775807
long long Ped_longlong = -123456789012345;
// unsigned long long range: 0 to 18446744073709551615
unsigned long long Ped_longlong_unsigned = 123456789012345;

// ----- printing -----
printf("char used for saving character: %c\n", Ped_RealChar);
printf("char used saving integer BUT the character is printed
      : %c\n", Ped_NumChar);
printf("char used saving integer: %d\n", Ped_NumChar);
printf("unsigned char saving integer: %u\n",
      Ped_NumChar_unsigned);
printf("short: %hd\n", Ped_sh);
printf("unsigned short: %hu\n", Ped_sh_unsigned);
printf("int: %d\n", Ped_int);

```

```
printf("unsigned int: %u\n", Ped_int_unsigned);
printf("long: %ld\n", Ped_long);
printf("unsigned long: %lu\n", Ped_long_unsigned);
printf("long long: %lld\n", Ped_longlong);
printf("unsigned long long: %llu\n", Ped_longlong_unsigned);
}
```

Now you have the source code `Pedram.c`, open it in VScode by executing `code Pedram.c` in the terminal. You should be able to see the code in the opened window. Now we need another terminal inside the VScode to compile and run the code. To open a terminal in VScode, go to the View menu at the top of the window. From the View menu, select "Terminal". You can see which directory this terminal is in by executing `pwd` in the terminal. Change your directory to the one where source code `Pedram.c` is. We need to do this so when we are compiling the code, the compiler can find the source code and translate it to the machine's language!

Let's checkout the code. A `char` is a type used to save a single character, like `a` or `G` or anything else. But you can assign a numeric value to a `char` variable in C. In fact, a `char` variable is internally represented as a small integer. So, you can assign a number within the range of -128 to 127 a `char` variable.

In this example, the decimal value 80 is assigned to the `char` variable `Ped_RealChar`, which I have picked this name to save this value. The `%c` format specifier in the `printf` statement is used to print the character representation of `Ped_RealChar`. In this case, it will print the character 'P', as the ASCII value 80 corresponds to the character 'P'.

So, while a `char` variable is primarily used to represent characters, it can also store numeric values within its valid range. We will learn more about characters and strings in section [Characters and strings](#). This is what the output should be:

```
char used for saving character: P
char used saving integer BUT the character is printed: P
char used saving integer: 80
unsigned char saving integer: 252
short: -1234
unsigned short: 56789
int: -42
unsigned int: 123456
long: -9876543210
unsigned long: 9876543210
long long: -123456789012345
```



```
unsigned long long: 123456789012345
```

Run this code after the lecture, by exceeding the limits and see the warnings **OR** errors **OR** **wrong** results. Let's say, change the value `Ped_NumChar_unsigned` to 256 which is higher than then maximum allowed for this type. The other thing you can do, is applying arithmetic operations on different types of variables that we have learned in section [Arithmetic Operations on Integers](#).

What is the problem with this code? Variable names are too long. I can just search for a variable in VScode to see the type of variable when it is initialized!

2.2.7 Fixed-width Integer types

In the section [Using `printf` and `limits.h` to Get the Limits of Integers](#) we talked about the limits of integer that might be different in different platforms. **How we can write a code that is portable to any OS?**

The C99 standard introduced fixed-width integer types in order to provide a consistent and portable way of specifying integer sizes across different platforms. Prior to C99, the sizes of integer types like `int` and `long` were implementation-dependent, which could lead to issues when writing code that relied on specific bit widths.

By adding fixed-width integer types such as `int8_t`, `int16_t`, `int32_t`, and `int64_t`, the C99 standard ensured that programmers had precise control over the sizes of their integer variables. These types have guaranteed widths in bits, making them useful in situations where exact bit-level manipulation with low-level systems is required.

To use the fixed-width integer types, the header `<stdint.h>` needs to be included. This header provides the type definitions for these fixed-width types, ensuring consistency across different platforms. By including `<stdint.h>`, programmers can use these types with confidence, knowing the exact size and range of the integers they are working with.

In addition to `<stdint.h>`, the header `inttypes.h` is included to access the placeholders associated with the fixed-width integer types. These format specifiers, such as `PRId8`, `PRId16`, and so on, enable proper printing and scanning of these types using the `printf` function. Make a new source code by `nano FixedInteger.c`, and paste the following code inside the file and save this program. Open the code in VScode using `code FixedInteger.c` and change the directory where the source code `FixedInteger.c` is in. Compile and execute the code by:

```
gcc -o FixedInteger FixedInteger.c
```

and

```
./FixedInteger.
```

```
#include <stdio.h>
#include <stdint.h>
#include <inttypes.h>

int main() {

    printf("Size of int8_t: %zu bits\n", 8 * sizeof(int8_t));
    printf("Signed int8_t range: %" PRIu8 " to %" PRIu8 "\n",
        INT8_MIN, INT8_MAX);
    printf("\n");

    printf("Size of uint8_t: %zu bits\n", 8 * sizeof(uint8_t));
    printf("uint8_t range: %d to %" PRIu8 "\n", 0, UINT8_MAX);
    printf("\n");

    printf("Size of int16_t: %zu bits\n", 8 * sizeof(int16_t));
    printf("Signed int16_t range: %" PRIu16 " to %" PRIu16 "\n",
        INT16_MIN, INT16_MAX);
    printf("\n");

    printf("Size of uint16_t: %zu bits\n", 8 * sizeof(uint16_t));
    printf("uint16_t range: %d to %" PRIu16 "\n", 0, UINT16_MAX);
    printf("\n");

    printf("Size of int32_t: %zu bits\n", 8 * sizeof(int32_t));
    printf("Signed int32_t range: %" PRIu32 " to %" PRIu32 "\n",
        INT32_MIN, INT32_MAX);
    printf("\n");

    printf("Size of uint32_t: %zu bits\n", 8 * sizeof(uint32_t));
    printf("uint32_t range: %d to %" PRIu32 "\n", 0, UINT32_MAX);
    printf("\n");

    printf("Size of int64_t: %zu bits\n", 8 * sizeof(int64_t));
    printf("Signed int64_t range: %" PRIu64 " to %" PRIu64 "\n",
        INT64_MIN, INT64_MAX);
    printf("\n");
```

```
printf("Size of uint64_t: %zu bits\n", 8 * sizeof(uint64_t));
printf("uint64_t range: %d to %" PRIu64 "\n", 0, UINT64_MAX);
}
```

The result not only in my machine, but also in any platform must be the same.

```
Size of int8_t: 8 bits
Signed int8_t range: -128 to 127

Size of uint8_t: 8 bits
uint8_t range: 0 to 255

Size of int16_t: 16 bits
Signed int16_t range: -32768 to 32767

Size of uint16_t: 16 bits
uint16_t range: 0 to 65535

Size of int32_t: 32 bits
Signed int32_t range: -2147483648 to 2147483647

Size of uint32_t: 32 bits
uint32_t range: 0 to 4294967295

Size of int64_t: 64 bits
Signed int64_t range: -9223372036854775808 to 9223372036854775807

Size of uint64_t: 64 bits
uint64_t range: 0 to 18446744073709551615
```

2.3 Characters and strings

In C, a string is defined as a sequence of characters. Individual characters are enclosed in single quotes `' '`, while strings are enclosed in double quotes `" "`.

To print a character, we use the placeholder `%c` in the `printf` function. For example, if `char c = 'P'`, then `printf("%c", c)`; will print the value of the character variable `c`.

To print a string, we use the placeholder `%s` in `printf`. For example, if `char s[] = "Pedram"`,

then `printf("%s", s);` will print the contents of the string variable `s`.

The size of a string can be determined using the `sizeof` operator (the same as finding the size of integer values), which returns the number of bytes occupied by the string. To print the size, we can use `%zu` as the placeholder in `printf`.

Strings in C are null-terminated, meaning they end with a **null character** (represented by 0). When accessing elements of a string, the index starts from 0, and the last character is always the null character. For example, in the `string s[] = "Pedram"`, `s[6]` refers to the null character. Don't forget the indexing in C starts from 0! Look at the following example:

```
#include <stdio.h>
// Compile and run the code with and without string.h
#include <string.h>

int main() {
    char c = 'P';
    char s[] = "Pedram";
    char s2[] = "Pasandide";

    printf("Character: %c\n", c);
    printf("String: %s\n", s);
    printf("s[0]: %c\n", s[0]);
    printf("s[5]: %c\n", s[5]);
    printf("Size of s: %zu\n", sizeof(s));
    printf("s[6] (null character): %d\n", s[6]);

    char s3[20]; // Make sure s3 has enough space to hold the
                 // concatenated string

    strcpy(s3, s); // Copy the content of s to s3
    strcat(s3, s2); // Concatenate s2 to s3

    printf("s3: %s\n", s3);

    return 0;
}
```

Open a terminal, checkout the directory you are in, using `pwd`. Make sure you are still in

```
/home/username/COMPSCI1XC3.
```

Make a new source code with `nano PedramString.c`, and copy-paste the code mentioned above. Open it in VScode with `code PedramString.c`. Change the directory to where you saved the source code. Compile the program with `gcc` and execute the code with `./PedramString`.

To concatenate two strings `s` and `s2` and store the result in `s3`, you can use the `strcpy` function from the `<string.h>` header.

In this code, `s` and `s2` are two strings that you want to concatenate. The variable `s3` is declared as an array of characters, with enough space to hold the concatenated string.

First, the `strcpy` function is used to copy the contents of `s` to `s3`, ensuring that `s3` initially holds the value of `s`. Then, the `strcat` function is used to concatenate `s2` to `s3`, effectively appending the contents of `s2` to `s3`.

Tips! There are many more functions dealing with string, and this one was just an example. Depending on your problem, you can search on Google and find how you can tackle your specific problem. Otherwise, remembering all these functions would be ALMOST impossible.

The result must be like:

```
Character: P
String: Pedram
s[0]: P
s[5]: m
Size of s: 7
s[6] (null character): 0
s3: PedramPasandide
```

2.4 Floating-point Numbers

In scientific programming, integers are often insufficient for several reasons:

- **1. Precision:** Integers have a finite range, and they cannot represent numbers with fractional parts. Many scientific computations involve non-integer values, such as real numbers, measurements, and physical quantities. Floating-point arithmetic allows for more precise representation and manipulation of these non-integer numbers.

- **2. Range:** While `long long int` can store larger integer values compared to regular `int`, it still has a limit. Scientific calculations often involve extremely large or small numbers, such as astronomical distances or subatomic particles. Floating-point numbers provide a wider range of values, accommodating these large and small magnitudes.

Floating-point numbers are represented and manipulated using floating-point arithmetic in CPUs. The encoding of floating-point numbers is based on the IEEE 754 standard, which defines formats for single-precision (32 bits) and double-precision (64 bits) floating-point numbers.

The basic structure of a floating-point number includes three components: the **sign** (positive or negative which can 0 or 1), the base (also known as the significant or **mantissa**), and the **exponent**. The base represents the significant digits of the number, and the exponent indicates the scale or magnitude of the number. Any floating-point number is represented in machine by:

$$(-1)^{\text{sign}} \times \text{mantissa} \times 2^{\text{exponent}}$$

For example, let's consider the number 85.3. In binary, it can be represented as approximately 1010101.01001100110011... In the IEEE 754 format, this number would be encoded as per the specifications of single-precision or double-precision floating-point representation. You can use online [converters](#) to get this number or you can read [more](#) how to do it. Right now in this course you don't need to necessarily learn how to do it, and I am mentioning this to illustrate everything clearly!

In the case of decimal fraction 0.3, its binary representation is non-terminating and recurring (0.0100110011...), meaning the binary fraction repeats infinitely. However, due to the finite representation of floating-point numbers in IEEE 754 format, the repeating binary fraction is rounded or truncated to fit the available number of bits. As a result, the exact decimal value of 0.3 cannot be represented accurately in binary using a finite number of bits.

So, when converting 0.3 to binary in the context of IEEE 754 floating-point representation, it will be approximated to the closest binary fraction that can be represented with the available number of bits. The accuracy of the approximation depends on the precision (number of bits) of the floating-point format being used.

The floating-point types used in C are `float`, `double` and `long double`. All these types are always signed meaning that they can represent both positive and negative value.

- `float` or **single-precision** has a precision of approximately 7 decimal digits. With smallest positive value of $1.17549435 \times 10^{-38}$ and largest positive value of $3.40282347 \times 10^{38}$
- `double` or **double-precision** has a precision of approximately 15 decimal digits. With smallest positive value of $2.2250738585072014 \times 10^{-308}$ and the largest positive value of $1.7976931348623157 \times 10^{308}$.

- `long double` or **extended-precision** format can vary in size depending on the platform. In x86 systems, it commonly uses 80 bits, but the specific number of bits for long double can differ across different architectures and compilers. In this course we won't use it!

Get the same results in **your machine** using the following code:

```
#include <stdio.h>
#include <float.h>

int main() {
    printf("Precision:\n");
    printf("Float: %d digits\n", FLT_DIG);
    printf("Double: %d digits\n", DBL_DIG);
    printf("Long Double: %d digits\n\n", LDBL_DIG);

    printf("Minimum and Maximum Values:\n");
    printf("Float: Minimum: %e, Maximum: %e\n", FLT_MIN, FLT_MAX);
    printf("Double: Minimum: %e, Maximum: %e\n", DBL_MIN, DBL_MAX);
    printf("Long Double: Minimum: %Le, Maximum: %Le\n", LDBL_MIN,
        LDBL_MAX);
}
```

The results for float and double in any computer should be the same showing the portability of these two types. I suggest you to use only these two at least in this course. In **my machine**, the results are:

```
Precision:
Float: 6 digits
Double: 15 digits
Long Double: 18 digits

Minimum and Maximum Values:
Float: Minimum: 1.175494e-38, Maximum: 3.402823e+38
Double: Minimum: 2.225074e-308, Maximum: 1.797693e+308
Long Double: Minimum: 3.362103e-4932, Maximum: 1.189731e+4932
```

Let's initialize a `double` value and print it using `printf`. This example defines a constant `double Ped` as 1.23456789 and demonstrates different printing options using the `%a.bf` placeholder, where `a` represents the minimum width and `b` specifies the number of digits after the decimal point:

```
#include <stdio.h>

int main() {
    const double Ped = 1.23456789;

    // Printing options with different placeholders
    // Minimum width = 0, 2 digits after decimal
    printf("Printing options for Ped = %.2f:\n", Ped);

    // Minimum width = 10, 4 digits after decimal
    printf("Printing options for Ped = %10.4f:\n", Ped);

    // Minimum width = 6, 8 digits after decimal
    printf("Printing options for Ped = %6.8f:\n", Ped);
}
```

In the first `printf` statement, `%.2f` is used, where 2 represents the minimum width (minimum number of characters to be printed) and .2 specifies 2 digits after the decimal point. This will print 1.23 as the output.

In the second `printf` statement, `%10.4f` is used. Here, 10 represents the minimum width, specifying that the output should be at least 10 characters wide, and .4 indicates 4 digits after the decimal point. This will print 1.2346 as the output, with 4 digits after the decimal and padded with leading spaces to reach a width of 10 characters.

In the third `printf` statement, `%6.8f` is used. The 6 represents the minimum width, and .8 specifies 8 digits after the decimal point. This will print 1.23456789 as the output, with all 8 digits after the decimal point.

By running this code, you can see the different printing options for the Ped value with varying widths and decimal precision.

```
Printing options for Ped = 1.23:
Printing options for Ped =      1.2346:
Printing options for Ped = 1.23456789:
```


2.4.1 Rounding Error in Floating-point Numbers

Rounding errors occur in floating-point arithmetic due to the finite number of bits allocated for representing the fractional part of a number. The rounding error becomes more prominent as we require higher precision or perform multiple arithmetic operations. The magnitude of the rounding error is typically on the order of the smallest representable number, which is commonly referred to as machine epsilon. **Run the following code!**

```
#include <stdio.h>

int main() {
    const float F = 1.23456789f;
    const double D = 1.23456789;
    const long double L = 1.23456789L;

    printf("Original values:\n");
    printf("Float: %.8f\n", F);
    printf("Double: %.8lf\n", D);
    printf("Long Double: %.8Lf\n\n", L);

    printf("Rounded values:\n");
    printf("Float: %.20f\n", F);
    printf("Double: %.20lf\n", D);
    printf("Long Double: %.20Lf\n", L);
}
```

In the provided code, the original values of `F`, `D`, and `L` are set to 1.23456789f, 1.23456789, and 1.23456789L, respectively. These values are printed with 8 digits of precision using `printf` statements.

When we examine the output, we can observe some differences between the original values and the rounded values. These differences arise due to the limitations of representing real numbers in the computer's finite memory using floating-point arithmetic.

In the original values section:

The original float value `F` is printed as 1.23456788, which differs from the original value due to the limited precision of the float data type. The original double value `D` is printed as 1.23456789, and in this case, there is no visible difference since the double data type provides sufficient precision to represent the value accurately. The original long double value `L` is also printed as 1.23456789,

indicating that the long double data type preserves the precision without any visible loss in this case.

In the rounded values section:

The float value `F` is printed with increased precision using `%.20f`. The rounded value is 1.23456788063049316406, which introduces rounding error due to the limited number of bits available for representing the fractional part of the number. The double value `D` is printed with increased precision using `%.201f`. Here, we can see a slight difference between the original value and the rounded value, with the rounded value being 1.23456788999999989009. This difference is attributed to the rounding error that occurs in the 16th digit after the decimal point. The long double value `L` is printed with increased precision using `%.20Lf`. The rounded value is 1.2345678900000000000003, demonstrating that even with the long double data type, there can still be a small rounding error.

In the case of `double` precision, the rounding error is approximately on the order of 10^{-16} , meaning that the least significant digit after the 16th decimal place can be subject to rounding error.

It's important to be aware of these limitations and potential rounding errors when performing calculations with floating-point numbers, especially in scientific and numerical computing, where high precision is often required.

```
Original values:
Float: 1.23456788
Double: 1.23456789
Long Double: 1.23456789

Rounded values:
Float: 1.23456788063049316406
Double: 1.23456788999999989009
Long Double: 1.2345678900000000000003
```

2.4.2 Type Conversion

In C, type conversion refers to the process of converting a value from one data type to another. There are two types of type conversion: explicit type conversion (also known as type casting) and implicit type conversion (also known as type coercion).

1. Implicit Type Conversion (Type Coercion): Implicit type conversion occurs automatically by the C compiler when performing operations between different data types. The conversion is done to ensure compatibility between operands. Here's an example:

```
#include <stdio.h>

int main() {
    int num = 10;
    double result = num / 3; // Implicitly convert int to double
    printf("Result: %f\n", result);

    int num2 = 10.6;
    printf("num2: %d\n", num2); // Implicitly convert double to int
}
```

In the above code, `num` is an integer with a value of `10`. When dividing it by `3`, the division operation requires a common data type for both operands. In this case, the compiler implicitly converts `num` to a double before performing the division. The result is stored in the `result` variable, which is of type double. When printing `result`, `%f` is used as the format specifier for floating-point numbers. In the other example, `num2` is also defined to be equal to `10.6`, but since the type is `int` anything after decimal will be neglected. Here is the output you should get:

```
Result: 3.000000
num2: 10
```

Implicit type conversion is performed based on a set of rules defined by the C language. It ensures that the operands are of compatible types to perform the desired operation.

2. Explicit Type Conversion (Type Casting): Explicit type conversion involves manually specifying the desired data type for the conversion. It is performed using type casting operators.

```
#include <stdio.h>

int main() {
    const int num1 = 10;
    const int num2 = 3;

    // 1. Print the division using int placeholder, ignoring
    // anything after the decimal
    int resultInt = num1 / num2;
    printf("Division without casting using int placeholder: %d\n",
        resultInt);

    // 2. Print the division using double placeholder without
    // casting (warning expected)
```

```

int resultDoubleNoCast = num1 / num2;
printf("Division without casting floating-point placeholder %f\n",
      resultDoubleNoCast);

// 3. Print the division by casting one of the operands
printf("Division (double with cast): %f\n", num1 / (double)num2);
;

// 4. Print the division by casting both operands
double resultDoubleBothCasted = (double)num1 / (double)num2;
printf("Division (double with both cast): %f\n",
      resultDoubleBothCasted);
}

```

Let's go through the code.

The division of `num1` by `num2` is stored in the `resultInt` variable, which is of type `int`. When using `%d` as the format specifier, the output will be an integer, ignoring anything after the decimal point.

The division without casting (`num1 num2`) is assigned to `resultDoubleNoCast`, which is of type `double`. However, this can lead to unexpected results due to integer division. The warning suggests that an implicit conversion is happening from `int` to `double`, which may not yield the desired precision.

To ensure proper division with decimal points, we cast one of the operands (`num2`) to `double` explicitly. This casting allows for a more accurate calculation. In some compiler you might see a warning here! To avoid that:

To perform the division correctly, both `num1` and `num2` are explicitly cast to `double`. This ensures that the division is carried out using floating-point arithmetic and produces the desired result. The `%f` format specifier is used to print the double value.

```

Division without casting using int placeholder: 3
Division without casting floating-point placeholder 0.000000
Division (double with cast): 3.333333
Division (double with both cast): 3.333333

```

2.4.3 Arithmetic Operations on Floating-point Numbers

The same as applying arithmetic operations on integer values, we can do the same on floating-point numbers. Here's a comprehensive example demonstrating arithmetic operations on floating-point numbers, including different data types, const and non-const values, and various arithmetic operators:

```
#include <stdio.h>

int main() {
    const float num1 = 10.5;
    float num2 = 5.2;
    double num3 = 7.8;
    const long double num4 = 3.14;
    long double num5 = 2.71;

    // Addition
    float sumFloat = num1 + num2;
    double sumDouble = num1 + num3;
    long double sumLongDouble = num4 + num5;

    printf("Addition:\n");
    printf("%.2f + %.2f = %.2f\n", num1, num2, sumFloat);
    printf("%.2f + %.2f = %.2f\n", num1, num3, sumDouble);
    printf("%.2Lf + %.2Lf = %.2Lf\n", num4, num5, sumLongDouble);
    printf("\n");

    // Subtraction
    float diffFloat = num1 - num2;
    double diffDouble = num1 - num3;
    long double diffLongDouble = num4 - num5;

    printf("Subtraction:\n");
    printf("%.2f - %.2f = %.2f\n", num1, num2, diffFloat);
    printf("%.2f - %.2f = %.2f\n", num1, num3, diffDouble);
    printf("%.2Lf - %.2Lf = %.2Lf\n", num4, num5, diffLongDouble);
    printf("\n");

    // Multiplication
    float productFloat = num1 * num2;
    double productDouble = num1 * num3;
    long double productLongDouble = num4 * num5;
```

```

printf("Multiplication:\n");
printf("%.2f * %.2f = %.2f\n", num1, num2, productFloat);
printf("%.2f * %.2f = %.2f\n", num1, num3, productDouble);
printf("%.2Lf * %.2Lf = %.2Lf\n", num4, num5, productLongDouble)
    ;
printf("\n");

// Division
float quotientFloat = num1 / num2;
double quotientDouble = num1 / num3;
long double quotientLongDouble = num4 / num5;

printf("Division:\n");
printf("%.2f / %.2f = %.2f\n", num1, num2, quotientFloat);
printf("%.2f / %.2f = %.2f\n", num1, num3, quotientDouble);
printf("%.2Lf / %.2Lf = %.2Lf\n", num4, num5, quotientLongDouble
    );
printf("\n");

// Compound Assignment Operators
num2 += num1;
num3 -= num4;
num5 *= num4;
float P1 = num1; // since I cannot change the value of num1
P1 /= num2;

printf("Compound Assignment Operators:\n");
printf("num2 += num1: %.2f\n", num2);
printf("num3 -= num4: %.2f\n", num3);
printf("num5 *= num4: %.2Lf\n", num5);
printf("num1 /= num2: %.2f\n", P1);
}

```

There are some compound assignment operators in the code including `+=`, `-=`, `*=`, and `/=`. How they work?

`+=` (Addition Assignment): It adds the value on the right-hand side to the variable on the left-hand side and assigns the result back to the variable. Example: `a += b`; is equivalent to `a =`

$a + b$;

-= (Subtraction Assignment): It subtracts the value on the right-hand side from the variable on the left-hand side and assigns the result back to the variable. Example: $a -= b$; is equivalent to $a = a - b$;

***=** (Multiplication Assignment): It multiplies the variable on the left-hand side by the value on the right-hand side and assigns the result back to the variable. Example: $a *= b$; is equivalent to $a = a * b$;

/= (Division Assignment): It divides the variable on the left-hand side by the value on the right-hand side and assigns the result back to the variable. Example: $a /= b$; is equivalent to $a = a / b$;

Note that we can use these compounds on inter values! The result you should get is:

Addition:

$10.50 + 5.20 = 15.70$

$10.50 + 7.80 = 18.30$

$3.14 + 2.71 = 5.85$

Subtraction:

$10.50 - 5.20 = 5.30$

$10.50 - 7.80 = 2.70$

$3.14 - 2.71 = 0.43$

Multiplication:

$10.50 * 5.20 = 54.60$

$10.50 * 7.80 = 81.90$

$3.14 * 2.71 = 8.51$

Division:

$10.50 / 5.20 = 2.02$

$10.50 / 7.80 = 1.35$

$3.14 / 2.71 = 1.16$

Compound Assignment Operators:

$\text{num2} += \text{num1}: 15.70$

$\text{num3} -= \text{num4}: 4.66$

$\text{num5} *= \text{num4}: 8.51$

$\text{num1} /= \text{num2}: 0.67$

2.5 Mathematical Functions

Mathematical functions in C are a set of functions provided by the standard library to perform common mathematical operations. These functions are declared in the `<math.h>` header file. They allow you to perform calculations involving numbers, trigonometry, logarithms, exponentiation, rounding, and more.

Some of the important and commonly used mathematical functions in C include:

- `sqrt(x)`: Calculates the square root of a number `x`.
- `pow(x, y)`: Raises `x` to the power of `y`.
- `fabs(x)`: Computes the absolute value of `x`.
- `sin(x)`, `cos(x)`, `tan(x)`: Computes the sine, cosine, and tangent of an angle `x`, respectively.
- `log(x)`: Computes the natural logarithm of `x`.
- `exp(x)`: Calculates the exponential value of `x`.
- `floor(x)`, `ceil(x)`, `round(x)`: Perform different types of rounding operations on `x`.
- `fmod(x, y)`: Calculates the remainder of dividing `x` by `y`.

Here is an example using these functions. Compile the code using `gcc -o pedram pedram.c`.

```
#include <stdio.h>
#include <math.h>

int main() {
    const double num1 = 4.0;
    const double num2 = 2.5;

    double sqrtResult = sqrt(num1);
    double powResult = pow(num1, num2);
    double sinResult = sin(num1);
    double logResult = log(num1);
    double ceilResult = ceil(num2);
    double fmodResult = fmod(num1, num2);

    printf("Square root of %.2f: %.2f\n", num1, sqrtResult);
}
```



```

printf("%.2f raised to the power of %.2f: %.2f\n", num1, num2,
    powResult);
printf("Sine of %.2f: %.2f\n", num1, sinResult);
printf("Natural logarithm of %.2f: %.2f\n", num1, logResult);
printf("Ceiling value of %.2f: %.2f\n", num2, ceilResult);
printf("Remainder of %.2f divided by %.2f: %.2f\n", num1, num2,
    fmodResult);
}

```

Probably, you see the following error:

```

/usr/bin/ld: /tmp/ccar3sn3.o: in function 'main':
pedram.c:(.text+0x30): undefined reference to 'sqrt'
/usr/bin/ld: pedram.c:(.text+0x50): undefined reference to 'pow'
/usr/bin/ld: pedram.c:(.text+0x67): undefined reference to 'sin'
/usr/bin/ld: pedram.c:(.text+0x7e): undefined reference to 'log'
/usr/bin/ld: pedram.c:(.text+0x95): undefined reference to 'ceil'
/usr/bin/ld: pedram.c:(.text+0xb5): undefined reference to 'fmod'
collect2: error: ld returned 1 exit status

```

Why? By default, the GCC compiler includes standard C libraries, but it does not automatically include all other libraries such as the `math` library. Therefore, when you use math functions like `sqrt`, `pow`, or `log`, the linker needs to know where to find the implementation of these functions.

Including the `<math.h>` header file in your code is necessary to provide the function prototypes and declarations for the math functions. It allows the compiler to understand the function names, parameter types, and return types when you use these functions in your code. However, including `<math.h>` alone is not sufficient to resolve references to the math functions during the linking phase (run the code without including `<math.h>` and check the errors!). The math library (`libm`) that contains **the actual implementation** of the math functions needs to be linked explicitly.

To solve this issue, the `-lm` flag explicitly tells the compiler to link with the math library (`libm`), allowing it to resolve the references to the math functions used in your code. You don't need to remember all these flags you can always use Google. Now compile the code using `gcc -o pedram pedram.c -lm`. Execute the object, and the result must be:

```

Square root of 4.00: 2.00
4.00 raised to the power of 2.50: 32.00
Sine of 4.00: -0.76

```

```
Natural logarithm of 4.00: 1.39
Ceiling value of 2.50: 3.00
Remainder of 4.00 divided by 2.50: 1.50
```

2.6 Statements

2.6.1 Comparison: `if` and `switch`

1. `if` statement:

The `if` statement in C is a conditional statement that allows you to perform different actions based on the evaluation of a condition. It follows a general format:

```
if (condition) {
    // Code to execute if the condition is true
}
```

What does it mean? The `condition` is an expression that evaluates to either **true** or **false**. If the condition is **true**, the code inside the if block will be **executed**. If the condition is **false**, the code inside the if block will be **skipped**. Let's say if **a** greater than **b** (`a > b`), execute the code inside the `if` block. Note that `condition` is replaced by `a > b`, and it is called **Comparison operators**. Comparison operators used in `if` statements:

- `==`: Checks if two values are equal.
- `!=`: Checks if two values are not equal.
- `<`: Checks if the left operand is less than the right operand.
- `>`: Checks if the left operand is greater than the right operand.
- `<=`: Checks if the left operand is less than or equal to the right operand.
- `>=`: Checks if the left operand is greater than or equal to the right operand.

There are also some **Logical operators** for combining conditions. Let's say **A** is the first `condition` and **B** is the second `condition`:

- `!A`: Logical **NOT** operator. It means if the condition **A** is not true, execute the statement in the `if` block. For example. in `!(a > b)`, execute the code if **a** is **NOT** greater than **b**.

- `A || B`: Logical **OR** operator. It evaluates to true if **either** operand **A** or **B** is true. So, one of these conditions at least must be true to execute the `if` block.
- `A && B`: Logical **AND** operator. It evaluates to true only if both operands **A** and **B** are true. It means both **A** and **B** conditions **MUST** be **true**.

Another way to include more complex `if` statement with multiple conditions is to use the general format using `if`, `else if`, and `else`:

```
if (condition1) {  
    // Code to execute if condition1 is true  
} else if (condition2) {  
    // Code to execute if condition2 is true  
} else {  
    // Code to execute if all previous conditions are false  
}
```

In this format, each condition is checked sequentially. If the first condition is true, the corresponding code block is executed. If it's false, the next condition is checked. If none of the conditions are true, the code block inside the `else` block is executed. Here is an example:

```
#include <stdio.h>  
  
int main() {  
    int a = 5;  
    int b = 10;  
  
    if (a == b) {  
        printf("a is equal to b\n");  
    } else if (a != b) {  
        printf("a is not equal to b\n");  
    } else if (a < b) {  
        printf("a is less than b\n");  
    } else if (a > b) {  
        printf("a is greater than b\n");  
    } else if (a <= b) {  
        printf("a is less than or equal to b\n");  
    } else if (a >= b) {  
        printf("a is greater than or equal to b\n");  
    } else {  
        printf("None of the conditions are true\n");  
    }  
}
```

```
}

int A = 1;
int B = 0;

if (!A) {
    printf("A is false\n");
}

if (A || B) {
    printf("At least one of A or B is true\n");
}

if (A && B) {
    printf("Both A and B are true\n");
}
}
```

You should see the following results:

```
a is not equal to b
At least one of A or B is true
```

A and **B** were supposed to be conditions why they are equal to 0 and 1? In C, the value **0** is considered **false**, and any **non-zero** value is considered **true**. Take a look at the following example:

```
#include <stdio.h>

int main() {
    int condition1 = 0;
    int condition2 = 1;

    if (condition1) {
        printf("Condition 1 is true\n"); //statement 1.1
    } else {
        printf("Condition 1 is false\n"); //statement 1.2
    }

    if (condition2) {
```

```
    printf("Condition 2 is true\n"); //statement 2.1
} else {
    printf("Condition 2 is false\n"); //statement 2.2
}
```

The result is given here. based on the results we can say **statement 1.2** and **statement 2.1** are executed and the rest is skipped.

```
Condition 1 is false
Condition 2 is true
```

In C, the `stdbool.h` header provides a set of definitions for Boolean data types and values. It introduces the `bool` type, which is specifically designed to represent **Boolean values**. The `bool` type can have two possible values: `true` and `false`. This header also defines the constants `true` and `false` as macro **constants**.

Using `stdbool.h` and the `bool` type can improve code readability and express the intent more clearly when dealing with **Boolean values**. It enhances code portability, as it ensures consistent Boolean semantics across different platforms and compilers. Here's an example that demonstrates the usage of `stdbool.h` and the `bool` type:

```
#include <stdio.h>
#include <stdbool.h>

int main() {
    bool condition1 = true;
    bool condition2 = false;

    if (condition1) {
        printf("Condition 1 is true\n"); // statement 1.1
    } else {
        printf("Condition 1 is false\n"); // statement 1.2
    }

    if (condition2) {
        printf("Condition 2 is true\n"); // statement 2.1
    } else {
        printf("Condition 2 is false\n"); // statement 2.2
    }
}
```

The result is given here. based on the results we can say **statement 1.1** and **statement 2.2** are executed and the rest is skipped. In this example, we include the `stdbool.h` header and declare two `bool` variables `condition1` and `condition2`. We assign `true` to `condition1` and `false` to `condition2`. We then use these variables in if statements to check their values and print corresponding messages.

```
Condition 1 is true
Condition 2 is false
```

By using `stdbool.h` and the `bool` type, the code becomes more self-explanatory, as it explicitly shows the usage of Boolean values. The constants `true` and `false` provide clarity in expressing the intent of conditions, making the code more readable and maintainable.

2. `switch` statement:

The switch statement in C is a control flow statement that allows you to select one of several execution paths based on the value of an expression. It provides an alternative to using multiple `if` statements when you have a series of conditions to check against a single variable.

The general format of the switch statement in C is as follows:

```
switch (expression) {
    case value1:
        // code to be executed if expression matches value1
        break;
    case value2:
        // code to be executed if expression matches value2
        break;
    case value3:
        // code to be executed if expression matches value3
        break;
    // more cases...
    default:
        // code to be executed if expression doesn't match any case
}
```

Warning! The `break` statement is used to exit the `switch` statement after executing the corresponding code block. Without the `break` statement, the execution would **fall through** to the next case, resulting in unintended behavior.

Here's an example that demonstrates the usage of the `switch` statement:

```
#include <stdio.h>

int main() {
    int choice;

    printf("Enter a number between 1 and 3: ");
    scanf("%d", &choice);

    switch (choice) {
        case 1:
            printf("You chose option 1.\n");
            break;
        case 2:
            printf("You chose option 2.\n");
            break;
        case 3:
            printf("You chose option 3.\n");
            break;
        default:
            printf("Invalid choice.\n");
    }
}
```

Here, we have used `scanf` function. The `scanf` function in C is used to read input from the standard input (usually the keyboard) and assign it to variables based on specified format specifiers. It allows you to accept user input during program execution, making your program more interactive. The general format of the `scanf` function is `scanf(format, argument_list);`, where the `format` parameter specifies the format of the expected input, while the `argument_list` contains the addresses of variables where the input will be stored. Here we have `scanf("%d", &choice)`, which means the given number will be saved in `choice` with format of `int` since we used placeholder `%d`.

In this example, the user is prompted to enter a number between 1 and 3. The input is stored in the variable `choice`. The `switch` statement is then used to check the value of `choice` and execute the corresponding code block. The output will be different based on the value you enter every time you execute the object file by `./pedram`.

If the user enters 1, the code inside the case 1 block is executed, printing "You chose option 1." If the user enters 2, the code inside the case 2 block is executed, printing "You chose option 2." If

the user enters 3, the code inside the case 3 block is executed, printing "You chose option 3."

If the user enters any other value, the code inside the default block is executed, printing "Invalid choice."

More example of `scanf()`! The `scanf` function scans the input stream and tries to match the input with the specified format. It skips whitespace characters by default and stops scanning when it encounters a character that doesn't match the format specifier. Here's an example to illustrate the usage of `scanf`:

```
#include <stdio.h>

int main() {
    int age;
    float height;

    printf("Enter your age: ");
    scanf("%d", &age);

    printf("Enter your height in meters: ");
    scanf("%f", &height);

    printf("You are %d years old and %.2f meters tall.\n", age,
        height);
}
```

In this example, the `scanf` function is used to read user input for age and height. The `%d` format specifier is used to read an integer value, and the `%f` format specifier is used to read a floating-point value. The `&` operator is used to obtain the address of the variables `age` and `height` for `scanf` to store the input values.

After the input is read, the program prints the values of `age` and `height` using `printf`. It's important to note that `scanf` requires correct format specifiers to match the input data type. Failure to use the correct format specifier can lead to unexpected behavior or errors in your program. Additionally, input validation and error handling are crucial when using `scanf` to ensure that the input is valid and the expected values are successfully read.

2.6.2 Loops and Iterations: `while`

The `while` statement in C is a control flow statement that allows you to repeatedly execute a block of code as long as a specified condition is true. It provides a way to create loops in your program.

The general format of the while statement in C is as follows:

```
while (condition) {  
    // code to be executed while the condition is true  
}
```

The `condition` is a **Boolean** expression that determines whether the loop should continue or terminate. **As long as the condition evaluates to true**, the code inside the loop will be executed repeatedly. If the condition becomes false, the loop will be exited, and the program will continue with the next statement after the loop. Here's an example that demonstrates the usage of the `while` statement:

```
#include <stdio.h>  
  
int main() {  
    int count = 1;  
  
    while (count <= 5) {  
        printf("Count: %d\n", count);  
        count++;  
    }  
}
```

In this example, the while loop is used to print the value of the count variable as long as it is less than or equal to 5. The count variable is **incremented** by 1 in each iteration using the `count++` statement.

The while loop continues to execute as long as the condition `count <= 5` is true. Once the count value becomes 6, the condition becomes **false**, and the loop is terminated. Here is the output in terminal.

```
Count: 1  
Count: 2  
Count: 3
```

Count: 4

Count: 5

Something that I forgot to tell you! Incrementing and decrementing are unary **operators** in C that are used to increase or decrease the value of a variable by a specific amount.

The **increment** operator `++` is used to **increment** the value of a variable by 1. It can be applied as a prefix (`++x`) or a postfix (`x++`) operator. When used as a prefix, the increment operation is performed before the value is used in an expression. When used as a postfix, the increment operation is performed after the value is used in an expression. Similarly, the **decrement** operator `--` is used to decrease the value of a variable by 1. It follows the same prefix and postfix notation as the increment operator. Here is an example:

```
#include <stdio.h>

int main() {
    int x = 5;

    printf("Original value: %d\n", x);

    printf("After x++: %d\n", x++); // Postfix increment
    printf("After ++x: %d\n", ++x); // Prefix increment
    printf("After x--: %d\n", x--); // Postfix decrement
    printf("After --x: %d\n", --x); // Prefix decrement
}
```

and the result is:

```
Original value: 5
After x++: 5
After ++x: 7
After x--: 7
After --x: 5
```

2.6.3 Loops and Iterations: `do`

The `do` statement is another type of loop in C that is similar to the `while` loop. The main difference is that the `do` loop executes the code block first and then checks the condition. This guarantees that the code inside the loop will be executed **at least once**, even if the condition is

initially false.

The general format of the do statement in C is as follows:

```
do {  
    // code to be executed  
} while (condition);
```

Here's an example to illustrate the usage of the `do` statement:

```
#include <stdio.h>  
  
int main() {  
    int count = 1;  
  
    do {  
        printf("Count: %d\n", count);  
        count++;  
    } while (count <= 5);  
}
```

In this example, the `do` loop is used to print the value of the `count` variable and increment it by 1. The loop continues to execute as long as the condition `count <= 5` is true. Since the initial value of `count` is 1, the loop body will be executed once, and then the condition is checked. If the condition is true, the loop will repeat, and if the condition is false, the loop will be exited. Here is the result:

```
Count: 1  
Count: 2  
Count: 3  
Count: 4  
Count: 5
```

2.6.4 Loops and Iterations: `for`

The `for` statement in C is a looping construct that allows you to execute a block of code repeatedly based on a specific condition. It is typically used when you know the exact number of iterations or when you need to perform a specific action for a fixed range of values. The general format of the `for` statement is as follows:

```
for (initialization; condition; increment/decrement) {  
    // Code to be executed in each iteration  
}
```

```
}
```

The **initialization** step is used to initialize the loop control variable before the loop starts. It is typically used to set an initial value.

The **condition** is evaluated before each iteration. If the condition is true, the loop body is executed; otherwise, the loop terminates.

The **increment** or **decrement** step is performed after each iteration and updates the loop control variable. It is used to control the termination condition of the loop.

Here's an example that demonstrates the usage of the **for** loop to print numbers from 5 to 0 using **decrement**:

```
#include <stdio.h>

int main() {
    for (int i = 5; i >= 0; i--) {
        printf("%d \n", i);
    }
}
```

In this example, the loop is initialized with **int i = 5**, the condition is **i >= 0**, and **i** is decremented by **i--** after each iteration. The loop iterates as long as the condition **i >= 0** is **true**. In each iteration, the value of **i** is printed using **printf**.

```
5
4
3
2
1
0
```

2.6.5 Nested loops

Nested loops in C are loops that are placed inside another loop. They allow you to perform repetitive tasks in a structured and organized manner when dealing with multiple dimensions or when you need to iterate over a combination of rows and columns. The general structure of nested loops is as follows:

```
for (outer initialization; outer condition; outer increment/
    decrement) {
```

```
// Code before the inner loop

for (inner initialization; inner condition; inner increment/
    decrement) {
    // Code inside the inner loop
}

// Code after the inner loop
}
```

Here's an example that demonstrates nested loops by printing out a matrix-like pattern based on rows and columns:

```
#include <stdio.h>

int main() {
    int rows = 5;
    int columns = 3;

    for (int i = 1; i <= rows; i++) {
        for (int j = 1; j <= columns; j++) {
            printf("(%d, %d) ", i, j);
        }
        printf("\n");
    }
}
```

In this example, we have an outer loop that iterates over the rows and an inner loop that iterates over the columns. The outer loop is controlled by the variable `i`, which represents the current row, while the inner loop is controlled by the variable `j`, which represents the current column.

The inner loop prints the coordinates `(row, column)` for each position in the matrix-like pattern. After printing the values for a row, we insert a newline character using `printf("\n")` to move to the next row.

```
(1, 1) (1, 2) (1, 3)
(2, 1) (2, 2) (2, 3)
(3, 1) (3, 2) (3, 3)
(4, 1) (4, 2) (4, 3)
```

```
(5, 1) (5, 2) (5, 3)
```

As you can see, the nested loops allow us to iterate over each row and column combination, printing out the corresponding coordinates in the pattern.

Nested loops are commonly used when working with multi-dimensional arrays, matrix operations, nested data structures, or any situation that requires iteration over multiple levels or dimensions. They provide a powerful tool for handling complex repetitive tasks in a structured manner.

2.6.6 Loops and Iterations: Controlling the Loop

Controlling loops in C involves using certain statements like `break`, `continue`, and `goto` to alter the flow of execution within the loop. These statements provide control over how and when the loop iterations are affected or terminated.

- `break` statement is used to immediately exit the loop, regardless of the loop condition. When encountered, the program flow continues to the next statement after the loop. It is typically used to prematurely terminate a loop based on a specific condition.
- `continue` statement is used to skip the current iteration of the loop and move to the next iteration. It allows you to bypass the remaining code in the current iteration and start the next iteration immediately. It is often used to skip certain iterations based on a condition.
- `goto` statement allows you to transfer the control of the program to a labelled statement elsewhere in the code. It is a powerful but potentially risky construct, as it can lead to complex and less readable code. It is generally advised to use `goto` sparingly and with caution.

Here's an example that demonstrates the usage of `break`, `continue`, and `goto` statements within a loop:

```
#include <stdio.h>

int main() {
    int i;

    // Example with break
    for (i = 1; i <= 10; i++) {
        if (i == 5) {
            break;
        }
    }
}
```

```
    }
    printf("%d ", i);
}
printf("\n");

// Example with continue
for (i = 1; i <= 10; i++) {
    if (i % 2 == 0) {
        continue;
    }
    printf("%d ", i);
}
printf("\n");

// Example with goto
for (i = 1; i <= 3; i++) {
    printf("Outer loop iteration: %d\n", i);
    for (int j = 1; j <= 3; j++) {
        printf("Inner loop iteration: %d\n", j);
        if (j == 2) {
            goto end;
        }
    }
}
end:
printf("Goto example finished.\n");
}
```

In this example, we have three loops that demonstrate the usage of `break`, `continue`, and `goto`:

The first loop uses `break` to exit the loop when `i` is equal to 5. As a result, the loop terminates prematurely and only prints the numbers from 1 to 4.

The second loop uses `continue` to skip printing even numbers. When `i` is divisible by 2, the loop skips the remaining code and moves to the next iteration. As a result, only odd numbers from 1 to 10 are printed.

The third loop demonstrates the usage of `goto`. In this case, when the inner loop reaches iteration 2, the `goto` statement is encountered, and the control jumps to the `end` label, bypassing

the remaining iterations of the inner and outer loops. Finally, the message "Goto example finished" is printed.

These control statements provide flexibility and allow you to alter the flow of execution within loops, making your code more efficient and concise in certain situations. However, it's important to use them judiciously and maintain code readability and clarity.

```
1 2 3 4
1 3 5 7 9
Outer loop iteration: 1
Inner loop iteration: 1
Inner loop iteration: 2
Goto example finished.
```

4th W/1

2.6.7 Variable Scope

2.7 Arrays

2.8 Functions

Constant vs. non-constant arguments

Forward declaration and mutual recursion

Global variables

4th W/2 + remained

3 Intermediate Topics in C

including:

debugging

MakeFiles

5th W/1

pointers

Dynamic memory allocation

5th W/2

4 Advanced Topics in C

including:

structs and typedefs

input and output reading a file or output a file

6th W/1 parallel computing

further topics

6th W/2

5 A Short Overview of ANN